CRM Sales Analysis – Azure End-to-End Data Engineering Project

1. Project Overview

This project implements a complete **end-to-end Azure data engineering and analytics pipeline** using a CRM Sales dataset. The objective was to ingest raw CRM data, clean and transform it using Apache Spark, store it in a structured data lake, expose it through Azure Synapse Analytics, and finally build interactive dashboards in **Power BI** for business analysis and reporting.

The project closely follows real-world data engineering best practices, including layered data lake design, data quality checks, SQL-based analytics, and BI consumption.

Project link: https://github.com/KasturiDisale/CRM-dataset-azure-data-engineering-project

---

2. Dataset Description

The CRM dataset consists of multiple CSV files representing different business entities: - **Accounts** – company-level information such as sector, revenue, employees, and location - **Products** – product catalog with pricing and series information - **Sales Transactions** – opportunities, deal stages, sales values, and dates - **Sales Team** – sales agents, managers, and regional offices

These datasets together enable comprehensive analysis of sales performance, customer contribution, and product revenue.

---

3. High-Level Architecture

```
Source CSV Files
      ↓
Azure Data Factory (Ingestion)
      ↓
Azure Data Lake Storage Gen2 (Raw / Transformed Layers)
      ↓
Azure Databricks (Spark Data Processing)
      ↓
Azure Data Lake Storage Gen2 (Curated Data)
      ↓
Azure Synapse Analytics (Serverless SQL + Lake Database)
```

```
      ↓
Power BI (Semantic Model & Dashboards)
```

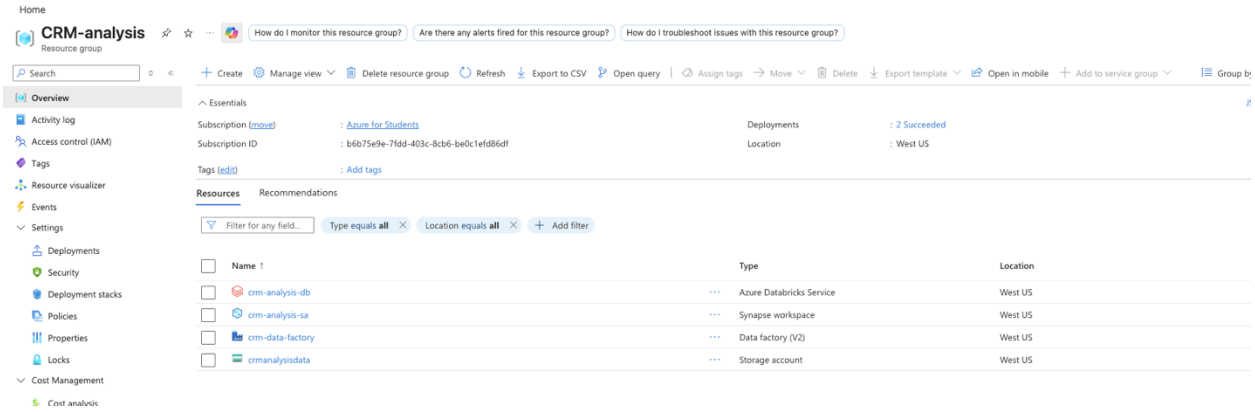

---

4. Azure Resource Setup

The following Azure services were created within a single resource group: - **Azure Data Factory** – for data ingestion pipelines - **Azure Data Lake Storage Gen2** – central data lake for raw and transformed data - **Azure Databricks** – Apache Spark environment for data processing - **Azure Synapse Analytics** – serverless SQL analytics and lake database - **Power BI** – reporting and visualization layer.

---

5. Data Ingestion using Azure Data Factory

Steps Performed

1. Created **Linked Services** for:
    - Source CRM CSV files
    - Azure Data Lake Storage Gen2
2. Built an ADF pipeline with multiple **Copy Data** activities
3. Ingested each dataset (accounts, products, sales transactions, sales team) into ADLS Gen2
4. Stored ingested files in a **raw-data** folder structure

Outcome

- All CRM source files were successfully ingested into the data lake
- Pipelines executed successfully and can be reused or scheduled

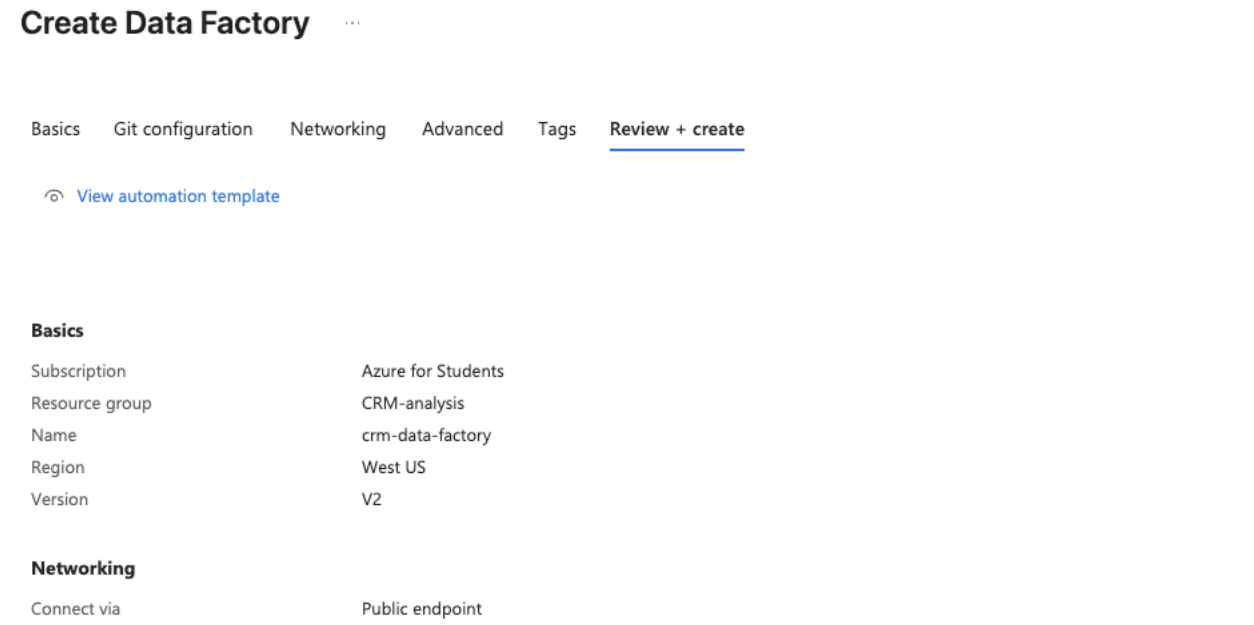

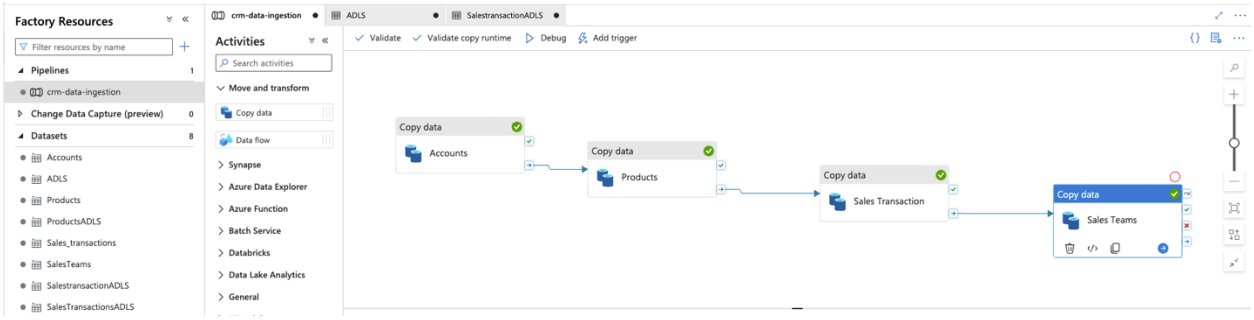

---

6. Data Lake Storage Design (ADLS Gen2)

The data lake was organized using a layered approach:

```
/raw-data
   ├── accounts
   ├── products
   ├── salestransaction
   └── salesteam

/transformed-data
   ├── accounts
```

```
├── products
├── salestransaction
└── salesteam
```

This design supports data traceability, reprocessing, and separation of raw and curated datasets.

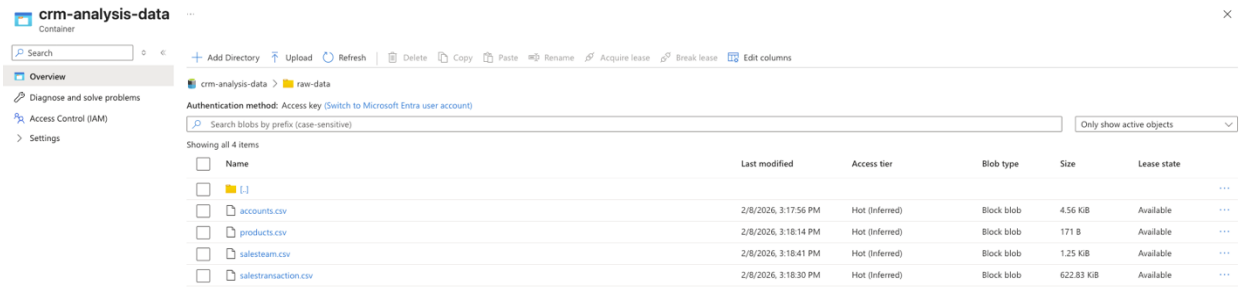

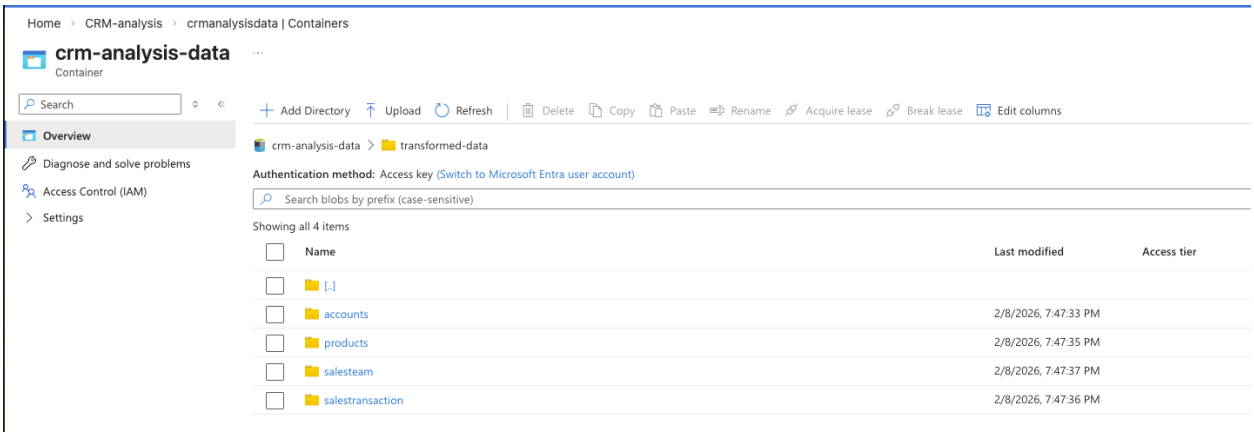

---

7. Data Processing using Azure Databricks (Apache Spark)

Azure Databricks was used to preprocess and clean the CRM data before analytics.

Key Processing Steps

1. Read raw CSV files from ADLS Gen2 using `abfss://` paths
2. Enabled schema inference and validated column data types
3. Performed **data quality checks**, including:
    - Column-wise null value analysis
    - Row count validation
4. Applied data cleaning logic:
    - Dropped records with null values in critical fields (e.g., account, close_value)
    - Retained or standardized optional fields where appropriate
5. Revalidated datasets to ensure no remaining nulls in required columns

6. Wrote cleaned and transformed data back to ADLS Gen2 in the **transformed-data** layer

Outcome

- High-quality, analytics-ready datasets were produced
- Data consistency and integrity were ensured before downstream consumption

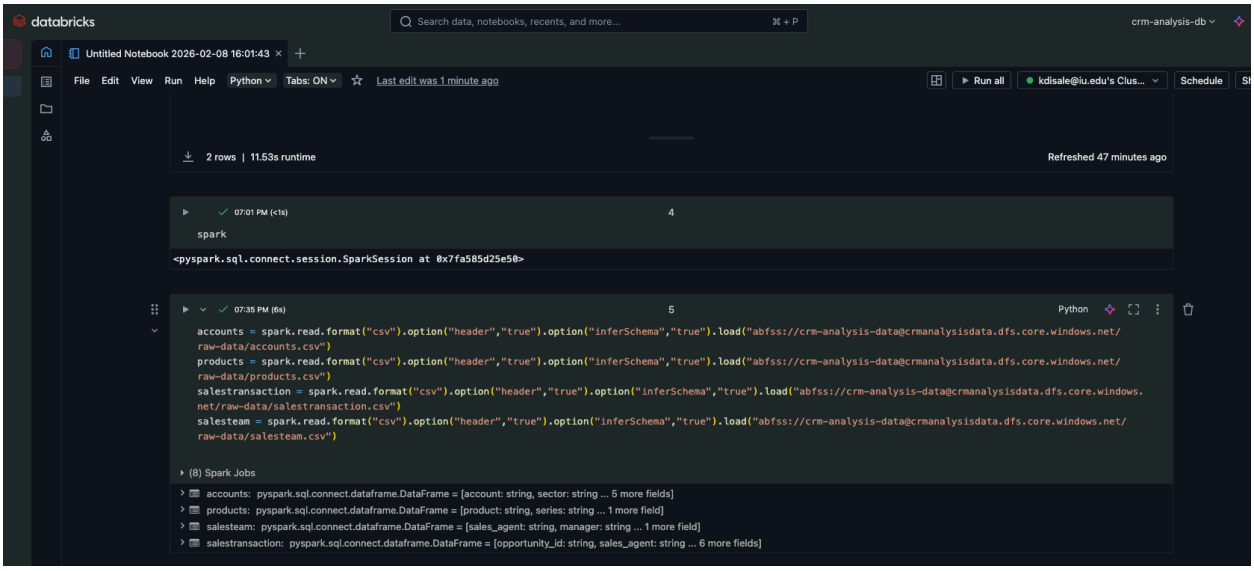

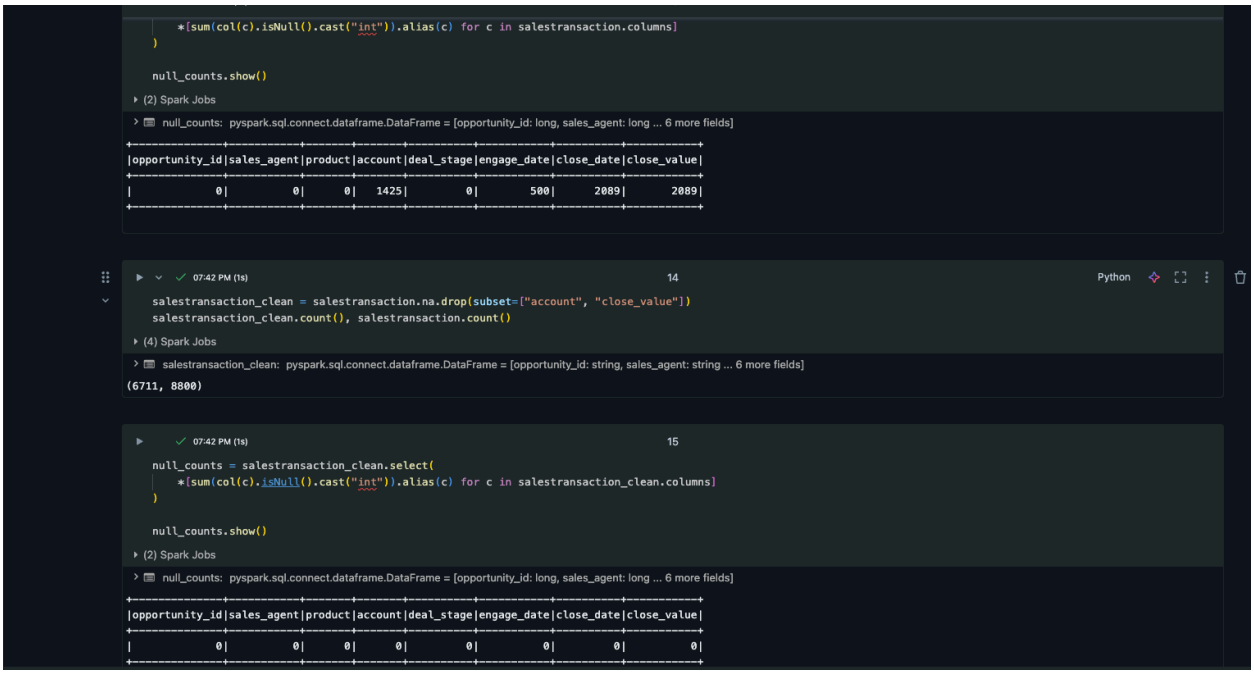

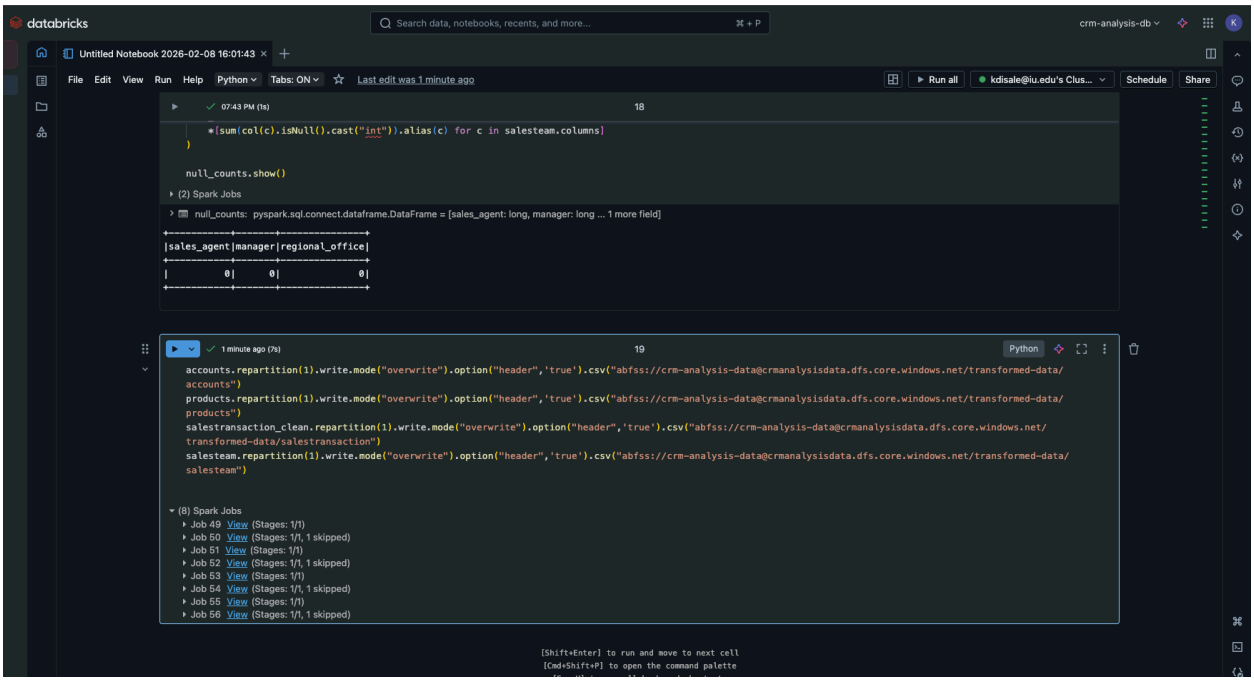

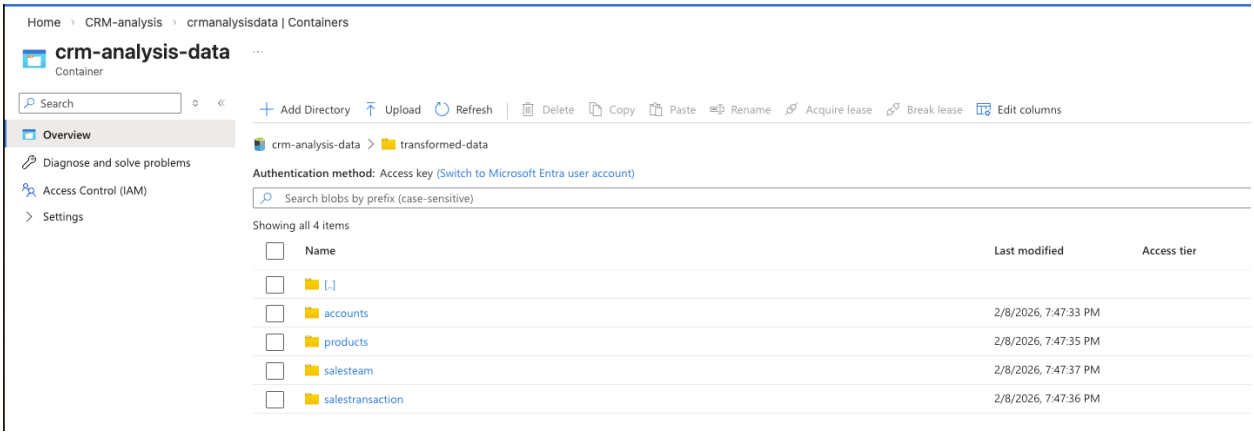

---

8. Azure Synapse Analytics – Lake Database & SQL

Synapse Setup

- Created a **Synapse Workspace** with the built-in **Serverless SQL Pool**
- Created a **Lake Database** named `CRManalysis_DB`

Tables Created

External tables were created over the transformed data for: - accounts - products - salestransaction - salesteam

These tables allow SQL-based querying directly on data lake files.

Analytics Performed

Using Synapse SQL, multiple analytical queries were executed, including: - Top 10 accounts by total revenue - Revenue by product and product series - Sales performance by sales agent - Deal counts and average deal sizes

Outcome

- Business-ready tables and queries exposed via SQL
- Serverless analytics enabled without infrastructure management

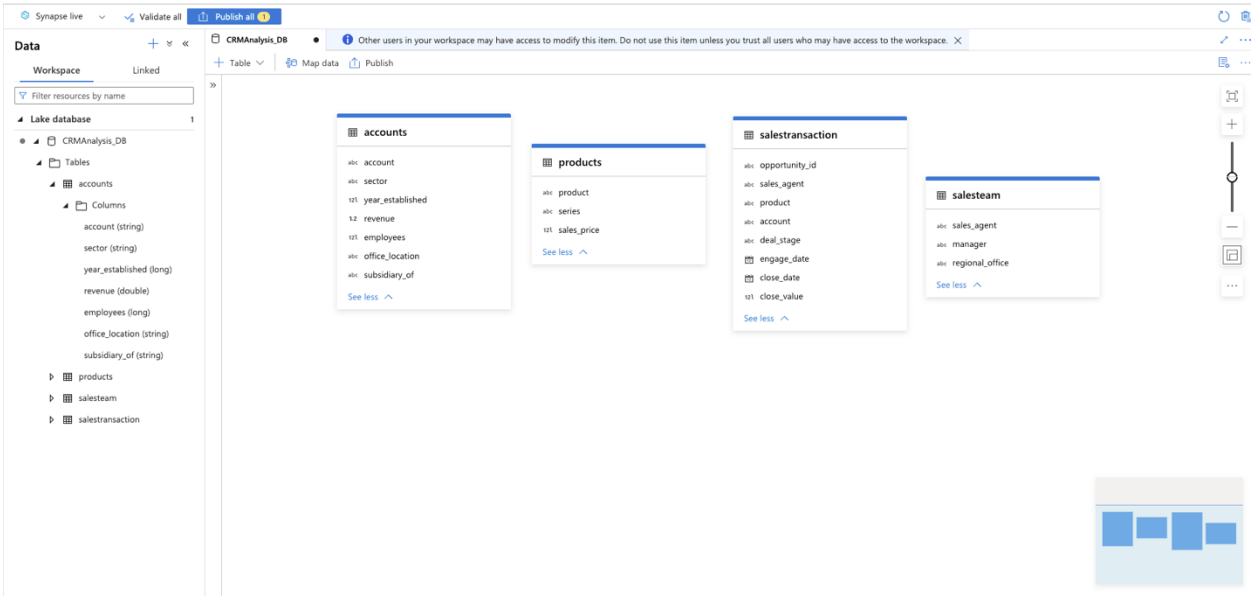

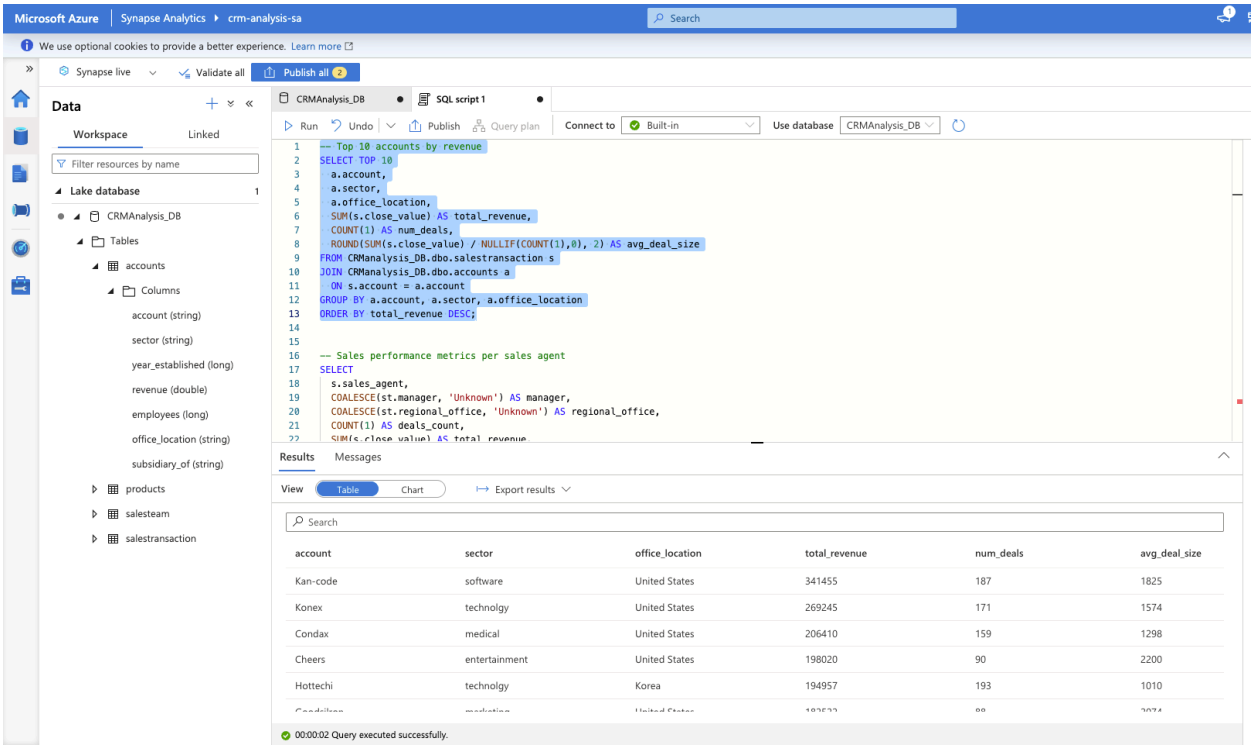
Microsoft Azure
Synapse Analytics
crm-analysis-sa
Search
We use optional cookies to provide a better experience. Learn more
Synapse live
Validate all
Publish all 2
Data
Workspace
Linked
Filter resources by name
Lake database 1
CRMAnalysis_DB
Tables
accounts
Columns
account (string)
sector (string)
year_established (long)
revenue (double)
employees (long)
office_location (string)
subsidiary_of (string)
products
salesteam
salestransaction
CRMAnalysis_DB
SQL script 1
Run
Undo
Publish
Query plan
Connect to
Built-in
Use database
CRMAnalysis_DB
-- Top 10 accounts by revenue
SELECT TOP 10
a.account,
a.sector,
a.office_location,
SUM(s.close_value) AS total_revenue,
COUNT(1) AS num_deals,
ROUND(SUM(s.close_value) / NULLIF(COUNT(1),0), 2) AS avg_deal_size
FROM CRMAnalysis_DB.dbo.salestransaction s
JOIN CRMAnalysis_DB.dbo.accounts a
ON s.account = a.account
GROUP BY a.account, a.sector, a.office_location
ORDER BY total_revenue DESC;
-- Sales performance metrics per sales agent
SELECT
s.sales_agent,
COALESCE(st.manager, 'Unknown') AS manager,
COALESCE(st.regional_office, 'Unknown') AS regional_office,
COUNT(1) AS deals_count,
Results
Messages
View
Table
Chart
Export results
Search
account
sector
office_location
total_revenue
num_deals
avg_deal_size
Kan-code
software
United States
341455
187
1825
Konex
technolgy
United States
269245
171
1574
Condax
medical
United States
206410
159
1298
Cheers
entertainment
United States
198020
90
2200
Hottechi
technolgy
Korea
194957
193
1010
00:00:02 Query executed successfully.

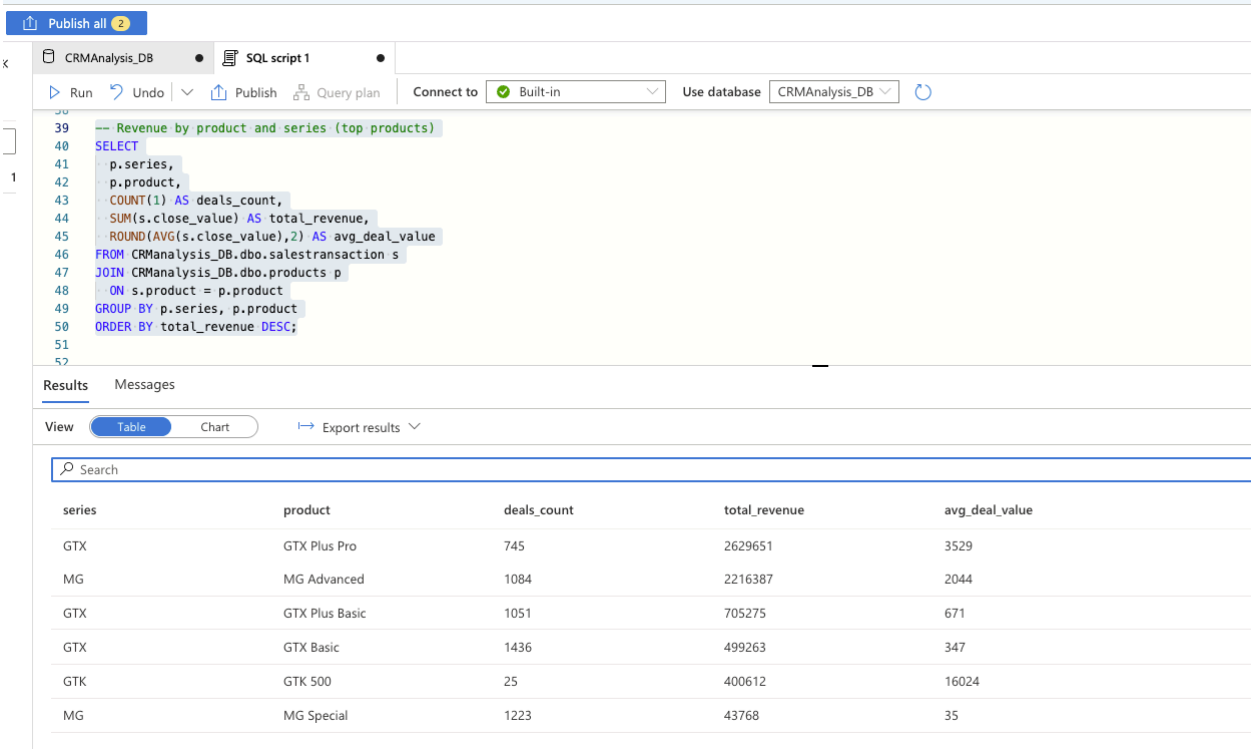
Publish all 2
CRMAnalysis_DB
SQL script 1
Run
Undo
Publish
Query plan
Connect to
Built-in
Use database
CRMAnalysis_DB
-- Revenue by product and series (top products)
SELECT
p.series,
p.product,
COUNT(1) AS deals_count,
SUM(s.close_value) AS total_revenue,
ROUND(AVG(s.close_value),2) AS avg_deal_value
FROM CRMAnalysis_DB.dbo.salestransaction s
JOIN CRMAnalysis_DB.dbo.products p
ON s.product = p.product
GROUP BY p.series, p.product
ORDER BY total_revenue DESC;
Results
Messages
View
Table
Chart
Export results
Search
series
product
deals_count
total_revenue
avg_deal_value
GTX
GTX Plus Pro
745
2629651
3529
MG
MG Advanced
1084
2216387
2044
GTX
GTX Plus Basic
1051
705275
671
GTX
GTX Basic
1436
499263
347
GTK
GTK 500
25
400612
16024
MG
MG Special
1223
43768
35

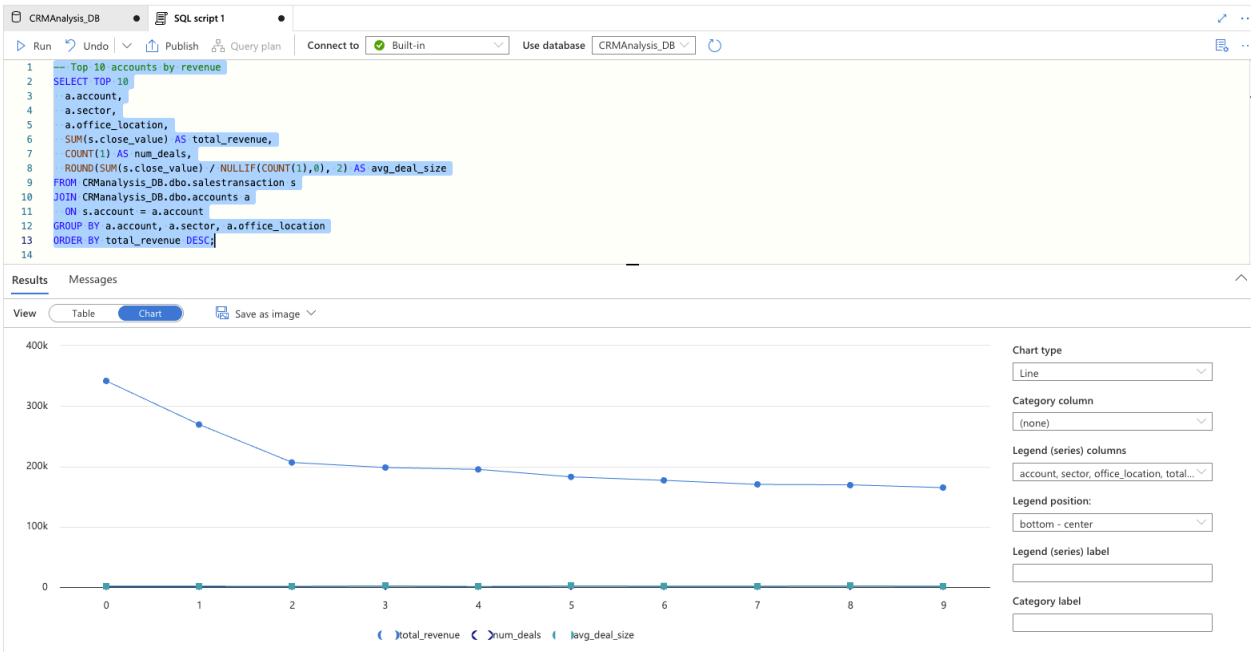

9. Power BI Integration & Dashboarding

Semantic Model

1. Connected Power BI to Synapse using the **Azure Synapse Analytics connector**
2. Created a **semantic model** from Synapse tables
3. Defined relationships in a **star schema**:
    - Fact table: salestransaction
    - Dimension tables: accounts, products, salesteam

Dashboards Built

The Power BI dashboard includes: - Average revenue by sector - Revenue by product and product series - Sales agent distribution by regional office - Interactive tables with account-level details - Slicers for sector, office location, and deal stage

Outcome

- Interactive and business-friendly dashboards
- End users can explore sales performance and trends dynamically

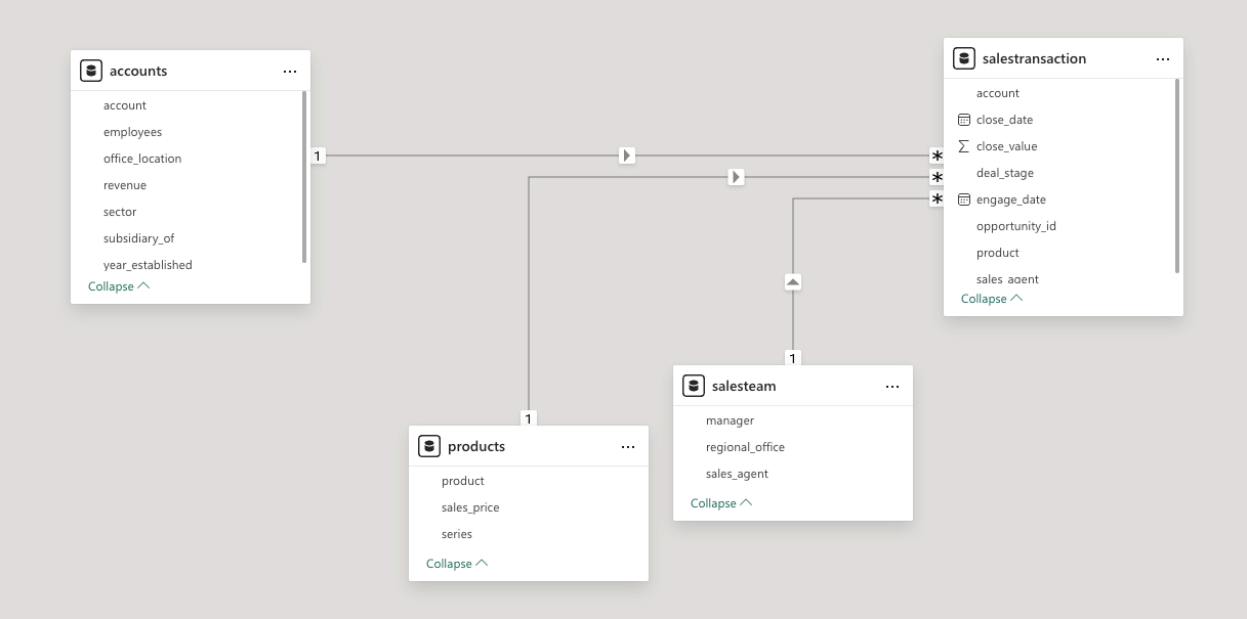

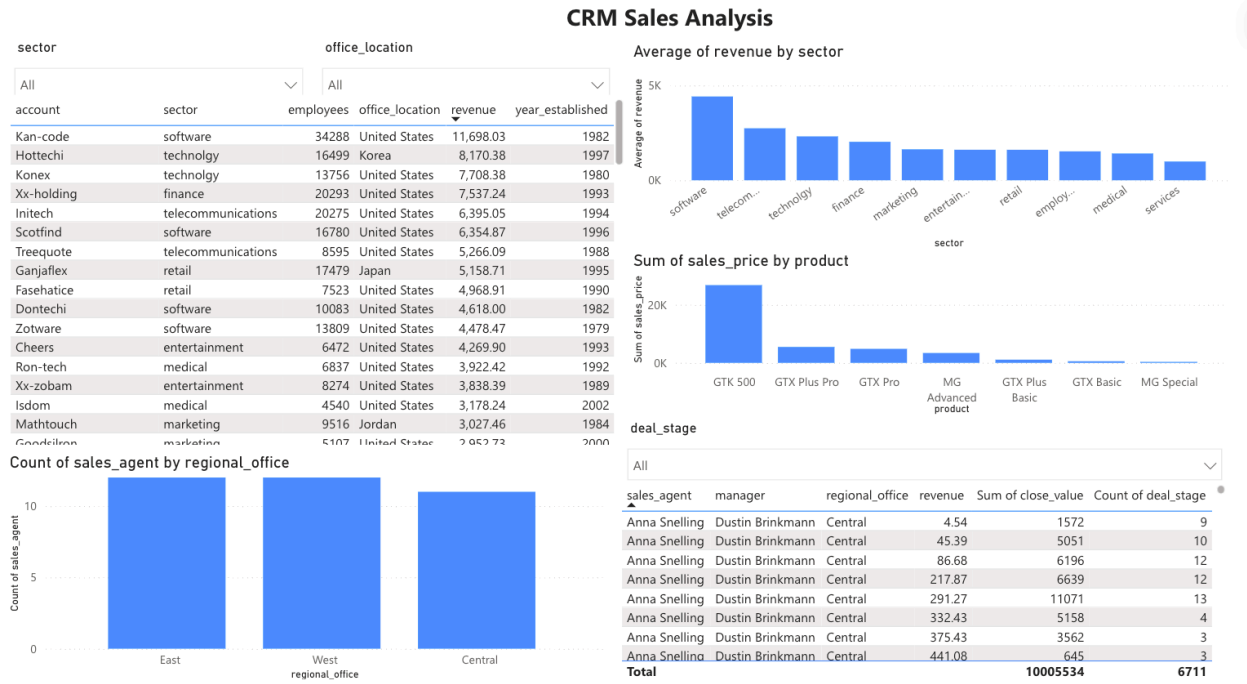

10. Tools & Technologies Used

- Azure Data Factory
- Azure Data Lake Storage Gen2
- Azure Databricks (Apache Spark, PySpark)
- Azure Synapse Analytics (Serverless SQL)
- Power BI

- SQL
- Python

---

11. Key Learnings

- Designing and implementing an end-to-end Azure data pipeline
- Applying data quality checks and transformations using Spark
- Using Synapse Serverless SQL for lakehouse analytics
- Building semantic models and dashboards in Power BI
- Applying star schema concepts for analytics

---

12. Conclusion

This project demonstrates a **real-world CRM sales analytics pipeline** built entirely on Azure. It showcases skills across data ingestion, big data processing, cloud analytics, and business intelligence. The architecture is scalable, cost-efficient, and aligned with industry best practices, making it suitable for portfolio demonstration and interview discussions.